

MONGODB

mongo design for blog posts

{
 _id
 author
 body

comments : {

date

permalink

tags : []

title

}

body
email
author



Strings

Supported data access patterns :



- > collecting most recent entries for home page
- > collecting all information to display one post
- > collecting all comments by a single author

bad access pattern for:
providing a table of contents by tag.

Alt. Schema for Blog

```
posts {
  _id ①
  title
  body
  author
  date
}
```

```
comments {
  _id
  post_id ①
  author
  author_email
  order: 0
}
```

```
tags {
  _id
  tag
  post_id ①
}
```

doesn't work particularly well

If it feels relational, it is probably wrong schema for MDB

Living w/o constraints

REL → foreign key constraints

MDB → data consistency is programmer's responsibility

> comments embedded in post → "pre-join"

→ embed data in ways that make sense

Life w/o transactions

MDB HAS ATOMIC OPERATIONS

— when working on a single doc.
no one else can access it until
you finished




approaches
to handle
trans.

one-to-one relations

employee: resume

building: floorplan



- frequency of access
 - if less frequent use ref (store separately)
- size : > 16mb must be REF.
 - anonymity of data

one-to-many
city: person

NYC
8m people

too large

duplication

~~city {
name
area
people: []
}~~

people {
name:
city: { name
area
...
}

TRUE LINKING

Two
collections

city {
-id: "NYC"
...
}

people {
name: "Jim"
city: NYC
...
}

One to few

↳ blog post : comments

posts {
 name
 comments : []
 ...
}

Many-to-many

BOOKS → AUTHORS

STUDENTS → TEACHERS

FEW : FEW (**USUALLY**)



* could embed books in authors doc

Students : teachers

teachers: [] students: []

MULTIKEYS (multikey index)

students

{ _id : 0
name: "Jim"
teachers: [0, 1, 2] }

teachers

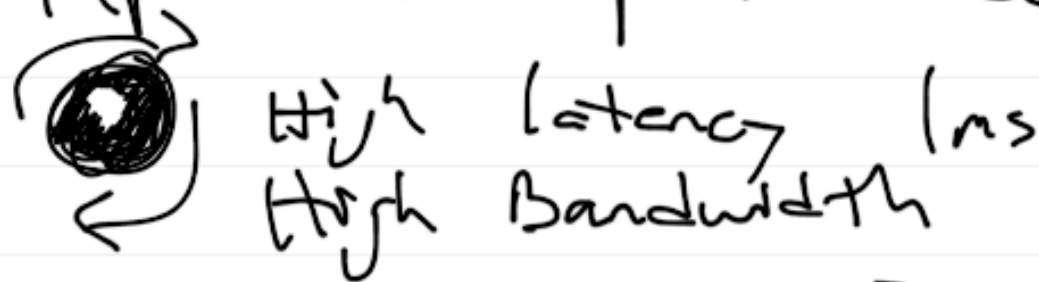
{ _id : 2
name: "TS" }

db.students.ensureIndex({ 'teachers': 1 })
db.students.find({ 'teachers':
{ \$all: [0, 1] } })

append .explain() to see what it did

Benefits of Embedding

- improved read performance



- one roundtrip to the DB

common categories

TREES

HOME : OUTDOORS : WINTER : SNOW

products:

category : 7
product_name :

category:

- id : 7

cat_name :

(1) [parent : 6]

options (2) children : []

(3) ancestors :

[3, 5, 9, 8]

(in order)

db.categories.find ({ ancestors : 34 })

- finds all documents when '34' is an element in the ancestors array

WHEN TO DENORMALIZE

1:1 — embedded

1:Many — (embed from the many to the 1)

many:many — link

to support applications
access pattern

embedding is perfectly acceptable

to go the other
way, link